

Kevin Gong
kg2445
HW01
Stat W4240
Section 2

Homework 1

Problem #1 (James 2.8)

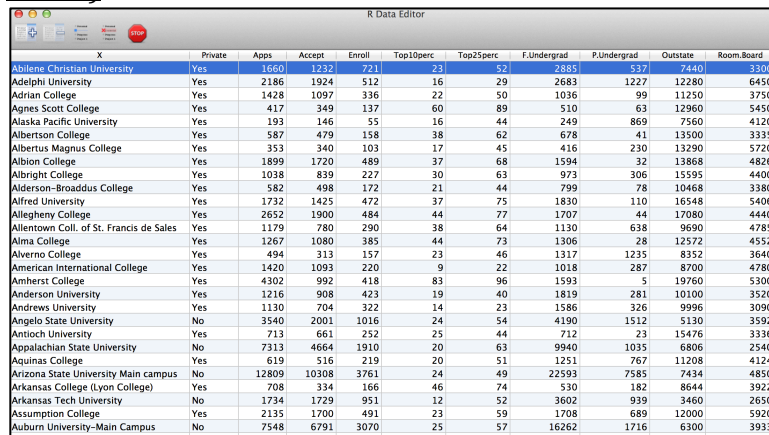
Part a)

```
> college <- read.csv('http://www-bcf.usc.edu/~gareth/ISL/College.csv')
> head(college)
```

		X	Private	Apps	Accept	Enroll	Top10perc	Top25perc	
1	Abilene Christian University	Yes	1660	1232	721	23	52		
2	Adelphi University	Yes	2186	1924	512	16	29		
3	Adrian College	Yes	1428	1097	336	22	50		
4	Agnes Scott College	Yes	417	349	137	60	89		
5	Alaska Pacific University	Yes	193	146	55	16	44		
6	Albertson College	Yes	587	479	158	38	62		
		F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	Personal	PhD	Terminal
1		2885	537	7440	3300	450	2200	70	78
2		2683	1227	12280	6450	750	1500	29	30
3		1036	99	11250	3750	400	1165	53	66
4		510	63	12960	5450	450	875	92	97
5		249	869	7560	4120	800	1500	76	72
6		678	41	13500	3335	500	675	67	73
		perc.alumni	Expend	Grad.Rate					
1		12	7041	60					
2		16	10527	56					
3		30	8735	54					
4		37	19016	59					
5		2	10922	15					
6		11	9727	55					

Here, we read the `College.csv` data into a variable we'll call `college`. While the question asks us to call the data `college` (and we've done this in our code), for the purposes of saving space here, we'll use `head(college)` to examine the first few rows of `college`.

Part b)



	X	Private	Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board
1	Abilene Christian University	Yes	1660	1232	721	23	52	2885	537	7440	3300
2	Adelphi University	Yes	2186	1924	512	16	29	2683	1227	12280	6450
3	Adrian College	Yes	1428	1097	336	22	50	1036	99	11250	3750
4	Agnes Scott College	Yes	417	349	137	60	89	510	63	12960	5450
5	Alaska Pacific University	Yes	193	146	55	16	44	249	869	7560	4120
6	Albertson College	Yes	587	479	158	38	62	678	41	13500	3335
7	Albertus Magnus College	Yes	353	340	103	17	45	416	230	13290	5720
8	Albion College	Yes	1899	1720	489	37	68	1594	32	13868	4826
9	Albright College	Yes	1038	839	227	30	63	973	306	15595	4400
10	Alderson-Broadus College	Yes	582	498	172	21	44	799	78	10468	3380
11	Alfred University	Yes	1732	1425	472	37	75	1830	110	16548	5406
12	Allegheny College	Yes	2652	1900	484	44	77	1707	44	17080	4440
13	Allentown Coll. of St. Francis de Sales	Yes	1179	780	290	38	64	1130	638	9690	4785
14	Alma College	Yes	1267	1080	385	44	73	1306	28	12572	4552
15	Alverno College	Yes	494	313	157	23	46	1317	1235	8352	3640
16	American International College	Yes	1420	1093	220	9	22	1018	287	8700	4780
17	Amherst College	Yes	4302	992	418	83	96	1593	5	19760	5300
18	Anderson University	Yes	1216	908	423	19	40	1819	281	10100	3520
19	Andrews University	Yes	1130	704	322	14	23	1586	326	9996	3090
20	Angelo State University	No	3540	2001	1016	24	54	4190	1512	5130	3592
21	Antioch University	Yes	713	661	252	25	44	712	23	15476	3336
22	Appalachian State University	No	7313	4664	1910	20	63	9940	1035	6806	2540
23	Aquinas College	Yes	619	516	219	20	51	1251	767	11208	4124
24	Arizona State University Main campus	No	12809	10308	3761	24	49	22593	7585	7434	4850
25	Arkansas College (Lyon College)	Yes	708	334	166	46	74	530	182	8644	3922
26	Arkansas Tech University	No	1734	1729	951	12	52	3602	939	3460	2650
27	Assumption College	Yes	2135	1700	491	23	59	1708	689	12000	5920
28	Auburn University-Main Campus	No	7548	6791	3070	25	57	16262	1716	6300	3933
29	Aurora College	No	443	433	147	13	30	1034	276	11063	4323

First we view the data using `fix(college)`

Using the provided code, we rename the row names and delete the first column of name data:

```
rownames(college) <- college[,1]
college <- college[,-1]
```

Now the first few columns look like this:

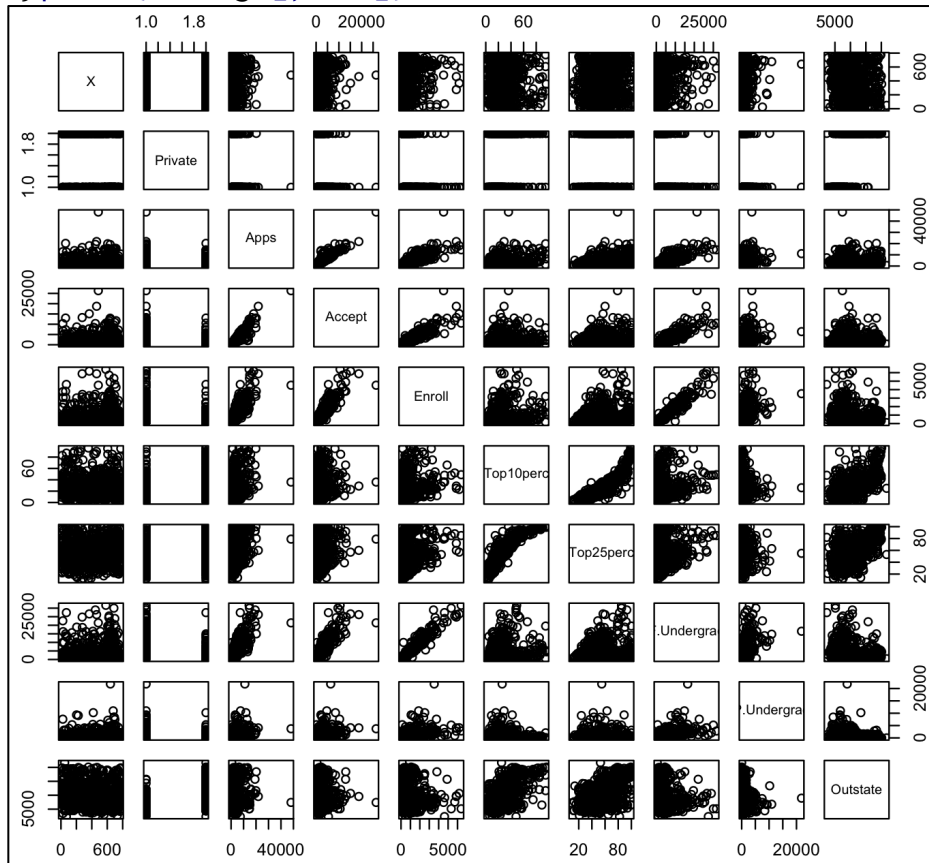
X	Private	Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad
Abilene Christian University	Yes	1660	1232	721	23	52	288
Adelphi University	Yes	2186	1924	512	16	29	268

Part c)


i) `summary(college)`

```
> summary(college)
      X      Private      Apps      Accept      Enroll      Top10perc      Top25perc      F.Undergrad      P.Undergrad      Outstate
Abilene Christian University: 1 No :212 Min. : 81 Min. : 72 Min. : 35 Min. : 1.00 Min. : 9.0 Min. : 139 Min. : 1.0 Min. : 2340
Adelphi University : 1 Yes:565 1st Qu.: 776 1st Qu.: 604 1st Qu.: 242 1st Qu.:15.00 1st Qu.: 41.0 1st Qu.: 992 1st Qu.: 95.0 1st Qu.: 7320
Adrian College : 1 Median :1558 Median :1110 Median : 434 Median :23.00 Median : 54.0 Median :1707 Median : 353.0 Median : 9990
Agnes Scott College : 1 Mean : 3002 Mean : 2019 Mean : 780 Mean :27.56 Mean : 55.8 Mean : 3700 Mean : 855.3 Mean :10441
Alaska Pacific University : 1 3rd Qu.: 3624 3rd Qu.: 2424 3rd Qu.: 902 3rd Qu.:35.00 3rd Qu.: 69.0 3rd Qu.: 4005 3rd Qu.: 967.0 3rd Qu.:12925
Albertson College : 1 Max. :48094 Max. :26330 Max. :6392 Max. :96.00 Max. :100.0 Max. :31643 Max. :21836.0 Max. :21700
(Other) :771
Room.Board      Books      Personal      PhD      Terminal      S.F.Ratio      perc.alumni      Expend      Grad.Rate
Min. :1780 Min. : 96.0 Min. : 850 Min. : 8.00 Min. : 24.0 Min. : 2.50 Min. : 0.00 Min. : 3186 Min. : 10.00
1st Qu.:3597 1st Qu.: 470.0 1st Qu.: 62.00 1st Qu.: 62.00 1st Qu.: 71.0 1st Qu.:11.50 1st Qu.:13.00 1st Qu.: 6751 1st Qu.: 53.00
Median :4200 Median : 500.0 Median :1200 Median : 75.00 Median : 82.0 Median :13.60 Median :21.00 Median : 8377 Median : 65.00
Mean :4358 Mean : 549.4 Mean :1341 Mean : 72.66 Mean : 79.7 Mean :14.09 Mean :22.74 Mean : 9660 Mean : 65.46
3rd Qu.:5050 3rd Qu.: 600.0 3rd Qu.:1700 3rd Qu.: 85.00 3rd Qu.: 92.0 3rd Qu.:16.50 3rd Qu.:31.00 3rd Qu.:10830 3rd Qu.: 78.00
Max. :8124 Max. :2340.0 Max. :6800 Max. :103.00 Max. :100.0 Max. :39.80 Max. :64.00 Max. :56233 Max. :118.00
```

ii) `pairs(college[,1:10])`



iii) `plot(college$Private,college$Outstate)`




We can see that private universities have significantly higher median out-of-state tuitions than public universities.

iv) `summary(college)`

```
Elite
Yes: 78
No  :699
```

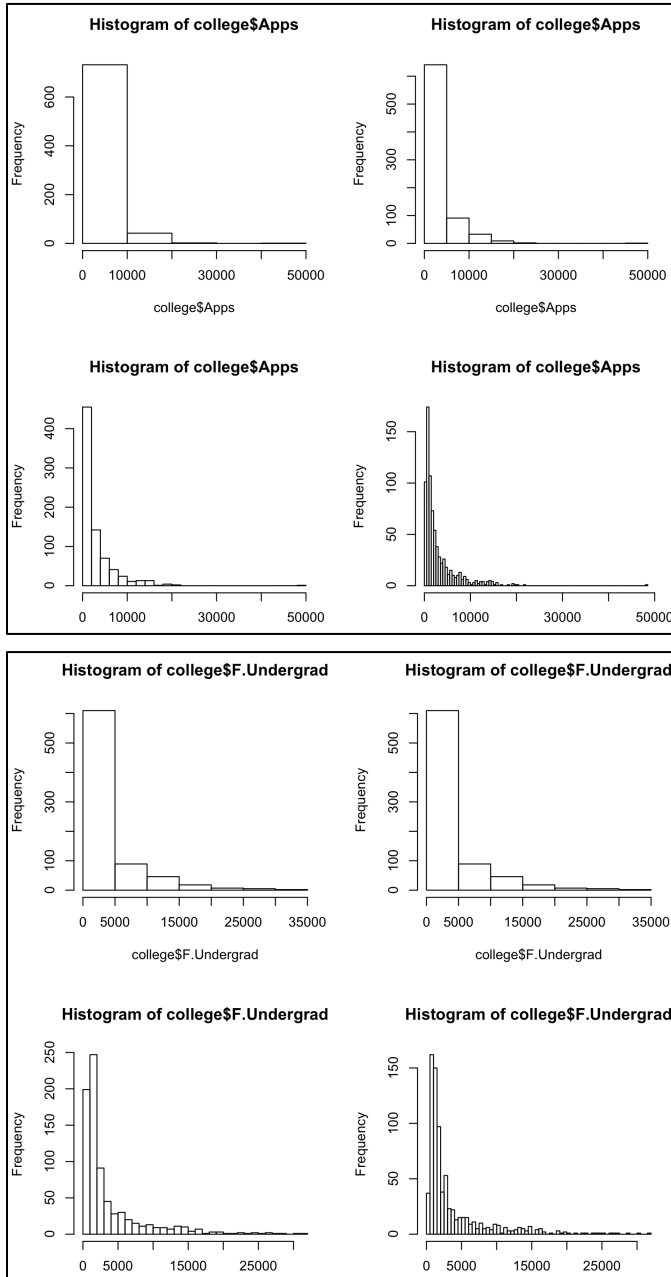
We see there are 78 elite universities, where an elite university is defined as one where the proportion of incoming students coming from the top 10% of their high school classes exceeds 50%.

`plot(college$Elite,college$Outstate)`

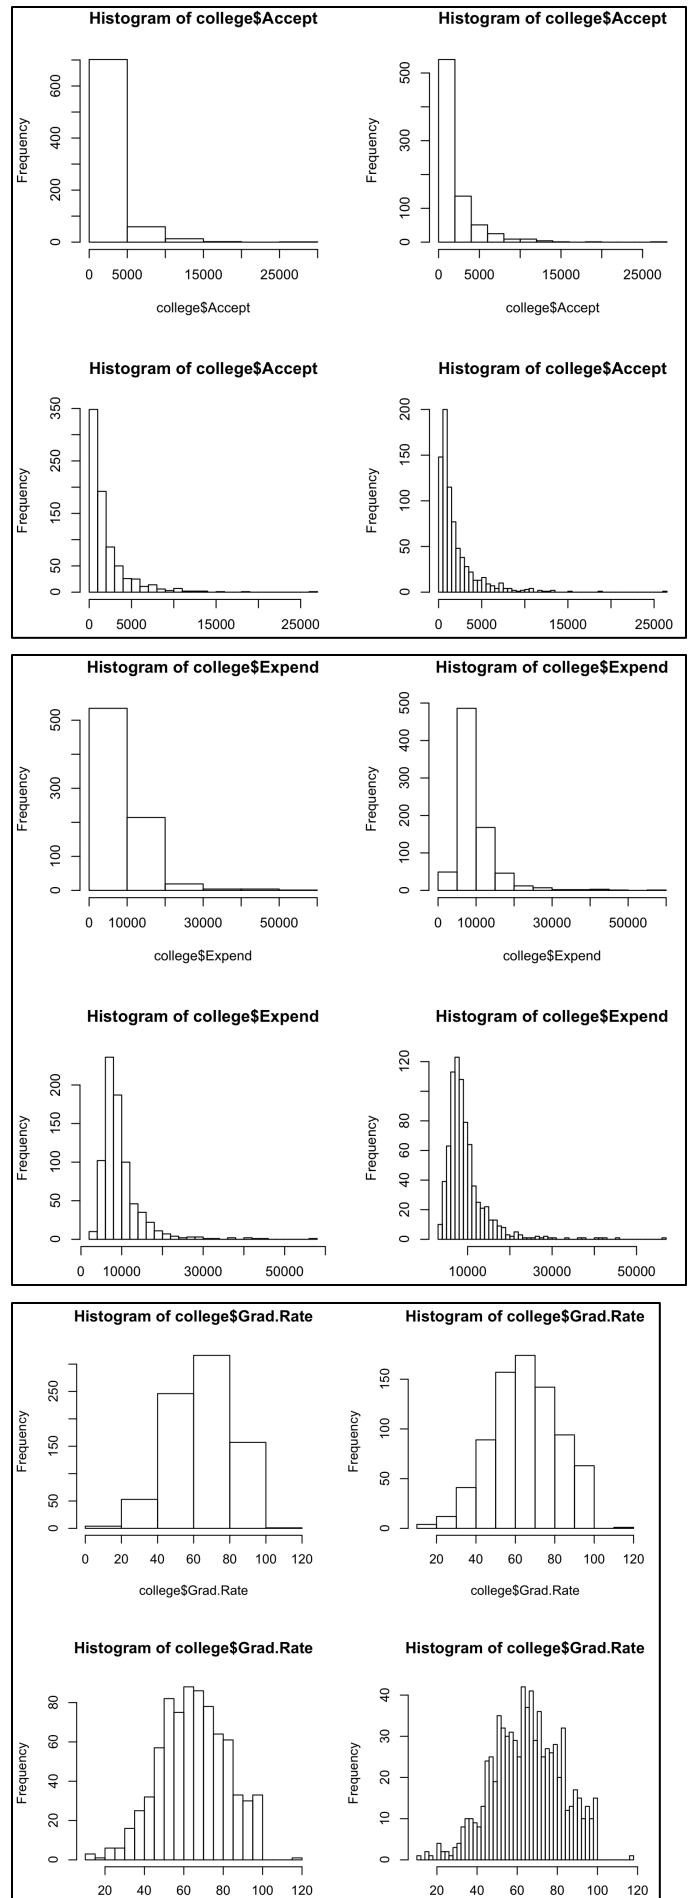


We can see that elite universities have significantly higher median out-of-state tuitions than non-elite universities.

v)



Histograms of # applications received, # applicants accepted, # full-time undergraduates, instructional expenditure per student, and graduate rate. The frequencies we used are 5, 10, 25, and 75.



vi) From the histograms in part v, we see that the mean graduation rate for all universities appears to lie between 60% and 80%, while the expense per student has a median of just under \$10,000 with a long right tail. From our boxplots, we can see that elite and private universities have significantly higher median out-of-state tuitions than non-elite universities.

Problem #2 (James 2.9)

Part a)

`sapply(auto,class)`

<code>sapply(auto,class)</code>	<code>mpg</code>	<code>cylinders</code>	<code>displacement</code>	<code>horsepower</code>	<code>weight</code>	<code>acceleration</code>	<code>year</code>	<code>origin</code>	<code>name</code>
	"numeric"	"integer"	"numeric"	"integer"	"integer"	"numeric"	"integer"	"integer"	"factor"

The *name* predictor is qualitative, while all other predictors are quantitative.

Part b)

`sapply(auto[,1:8],range,na.rm=TRUE)`

<code>> sapply(auto[,1:8],range,na.rm=TRUE)</code>	<code>mpg</code>	<code>cylinders</code>	<code>displacement</code>	<code>horsepower</code>	<code>weight</code>	<code>acceleration</code>	<code>year</code>	<code>origin</code>
[1,]	9.0	3	68	46	1613	8.0	70	1
[2,]	46.6	8	455	230	5140	24.8	82	3

The ranges for the quantitative predictors is shown above. This was done by applying the *range()* function to the first 8 predictors.

Part c)

`sapply(auto[,1:8],mean,na.rm=TRUE)`

`sapply(auto[,1:8],sd,na.rm=TRUE)`

<code>> sapply(auto[,1:8],mean,na.rm=TRUE)</code>	<code>mpg</code>	<code>cylinders</code>	<code>displacement</code>	<code>horsepower</code>	<code>weight</code>	<code>acceleration</code>	<code>year</code>	<code>origin</code>
	23.515869	5.458438	193.532746	104.469388	2970.261965	15.555668	75.994962	1.574307
<code>> sapply(auto[,1:8],sd,na.rm=TRUE)</code>	<code>mpg</code>	<code>cylinders</code>	<code>displacement</code>	<code>horsepower</code>	<code>weight</code>	<code>acceleration</code>	<code>year</code>	<code>origin</code>
	7.8258039	1.7015770	104.3795833	38.4911599	847.9041195	2.7499953	3.6900049	0.8025495

The mean and standard deviation of each quantitative predictor is shown above. This was done by applying the *mean()* and *sd()* functions to the quantitative predictors.

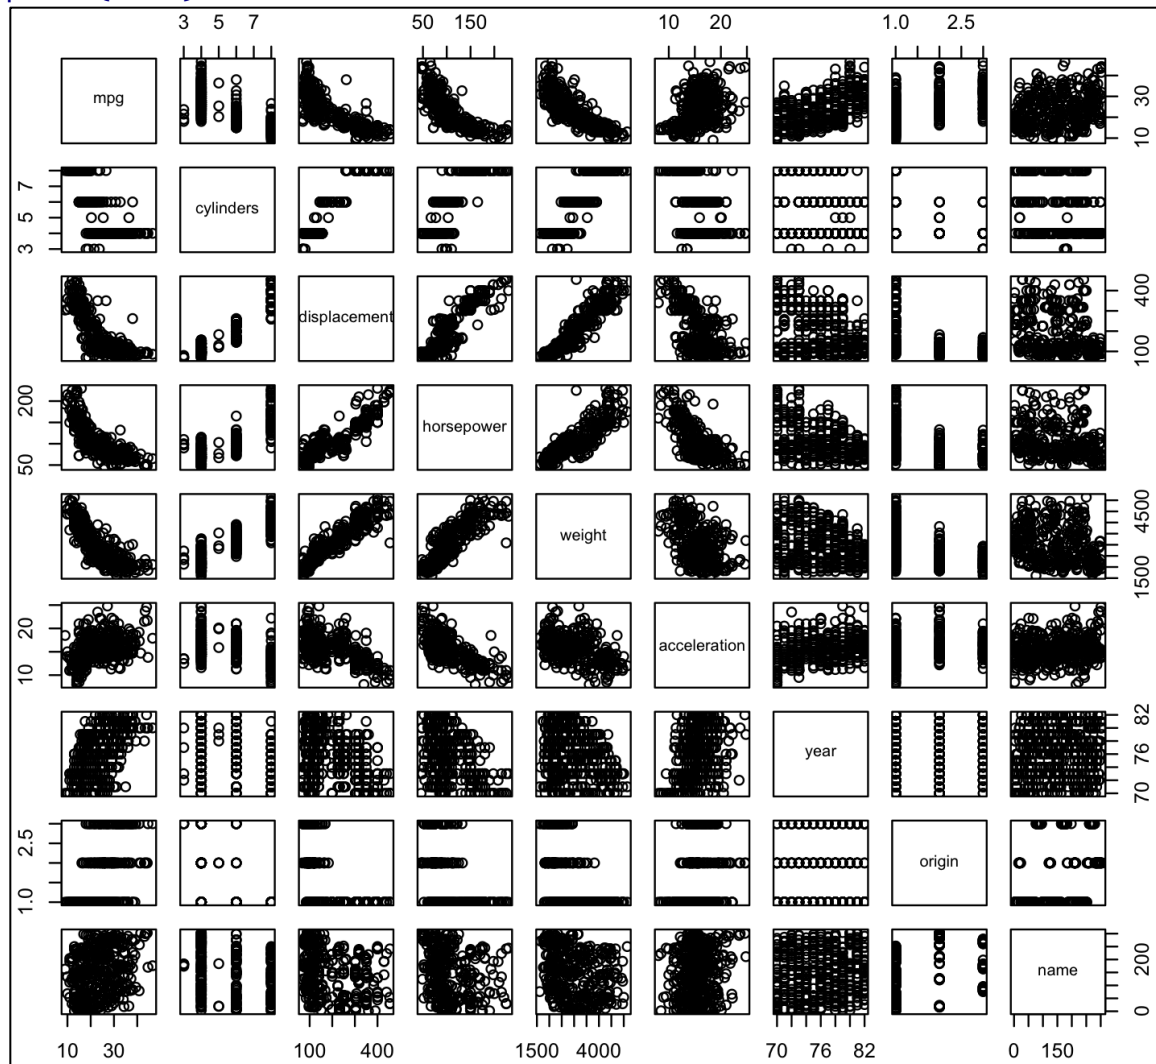
Part d)

```
auto = auto[-c(10:85),]
sapply(auto[,1:8],range,na.rm=TRUE)
sapply(auto[,1:8],mean,na.rm=TRUE)
sapply(auto[,1:8],sd,na.rm=TRUE)
```

```
> sapply(auto[,1:8],range,na.rm=TRUE)
      mpg cylinders displacement horsepower weight acceleration year origin
[1,] 11.0         3          68         46  1649          8.5  70      1
[2,] 46.6         8         455        230  4997        24.8  82      3
> sapply(auto[,1:8],mean,na.rm=TRUE)
      mpg cylinders displacement horsepower weight acceleration year origin
24.438629  5.370717 187.049844 100.955836 2933.962617 15.723053 77.152648 1.598131
> sapply(auto[,1:8],sd,na.rm=TRUE)
      mpg cylinders displacement horsepower weight acceleration year origin
7.9081842  1.6534857 99.6353853 35.8955668 810.6429384  2.6805138 3.1112298 0.8161627
```

Part e)

`pairs(auto)`



Part f)

From our scatterplot matrix in part e, we can see that *mpg* appears to have somewhat negative linear relationships with the predictors *displacement*, *horsepower*, and *weight*. The negative relationships indicate that mpg is higher when a car's displacement, horsepower, and weight are lower. Thus, these three

predictors could be useful in predicting *mpg* since they appear have fairly strong negative relationships with *mpg*.

Problem #3 (James 2.10)

Part a)


?Boston

Description

The `Boston` data frame has 506 rows and 14 columns.

The rows represent different tracts/towns/suburbs of Boston, while the columns represent different predictors/indicators.

Part b)



From a quick glance at the scatterplot matrix, there don't appear to be any especially strong linear correlations between the predictors that immediately stand out. Perhaps the most apparent ones are *rm* (average number of rooms per dwelling) correlating negatively with *lstat* (lower status of the population as a percent), *rm* correlating positively with *medv* (median value of owner-occupied homes), and a negative correlation between *lstat* and *medv*. These correlations seem to make sense since larger homes (which have more rooms than smaller homes) are probably higher in value and lower-status families are less likely to be able to afford these homes.

Part c)

The predictor for per capita crime rate is *crim*. From the scatterplot matrix, some more notable relationships suggest that:

- crime is higher among tracts that bound the Charles River than those that don't (*chas*)
- crime is higher where there are above average concentrations of nitrogen oxides (*nox*)
- crime appears to be exponentially correlated with the proportion of owner-occupied units built prior to 1940 (*age*)
- crime appears to be negatively exponentially correlated with the weighted mean of distances to five Boston employment centers (*dis*)
- crime appears to be negatively exponentially correlated with the median value of owner-occupied homes (*medv*)

Part d)

```
sapply(Boston, range, na.rm=TRUE)
```

```
sapply(Boston, mean, na.rm=TRUE)
```

```
sapply(Boston, sd, na.rm=TRUE)
```

```
summary(Boston)
```

```
> sapply(Boston, range, na.rm=TRUE)
      crim zn indus chas  nox   rm   age   dis   rad   tax ptratio black lstat medv
[1,] 0.00632 0 0.46  0 0.385 3.561 2.9 1.1296 1 187 12.6 0.32 1.73 5
[2,] 88.97620 100 27.74 1 0.871 8.780 100.0 12.1265 24 711 22.0 396.90 37.97 50
> sapply(Boston, mean, na.rm=TRUE)
      crim zn indus chas  nox   rm   age   dis   rad   tax ptratio black lstat medv
22.53280632
> sapply(Boston, sd, na.rm=TRUE)
      crim zn indus chas  nox   rm   age   dis   rad   tax ptratio black lstat medv
8.6015451 23.3224530 6.8603529 0.2539940 0.1158777 0.7026171 28.1488614 2.1057101 8.7072594 168.5371161 2.1649455 91.2948644 7.1410615 9.1971041
> summary(Boston)
      crim      zn      indus      chas      nox      rm      age      dis      rad      tax      ptratio      black      lstat      medv
Min.   :0.00632 Min.   :0.00 Min.   :0.46 Min.   :0.00000 Min.   :0.3850 Min.   :3.561 Min.   :2.90 Min.   :1.130 Min.   :1.000 Min.   :187.0
1st Qu.:0.08204 1st Qu.:0.00 1st Qu.:5.19 1st Qu.:0.00000 1st Qu.:0.4490 1st Qu.:5.886 1st Qu.:45.02 1st Qu.:2.100 1st Qu.:4.000 1st Qu.:279.0
Median :0.25651 Median :0.00 Median :9.69 Median :0.00000 Median :0.5380 Median :6.208 Median :77.50 Median :3.207 Median :5.000 Median :330.0
Mean   :3.61352 Mean   :11.36 Mean :11.14 Mean :0.06917 Mean :0.5547 Mean :6.285 Mean :68.57 Mean :3.795 Mean :9.549 Mean :408.2
3rd Qu.:3.67708 3rd Qu.:12.50 3rd Qu.:18.10 3rd Qu.:0.00000 3rd Qu.:0.6240 3rd Qu.:6.623 3rd Qu.:94.08 3rd Qu.:5.188 3rd Qu.:24.000 3rd Qu.:666.0
Max.   :88.97620 Max.   :100.00 Max. :27.74 Max. :1.00000 Max. :0.8710 Max. :8.780 Max. :100.00 Max. :12.127 Max. :24.000 Max. :711.0
      ptratio      black      lstat      medv
Min.   :12.60 Min.   :0.32 Min.   :1.73 Min.   :5.00
1st Qu.:17.40 1st Qu.:375.38 1st Qu.:6.95 1st Qu.:17.02
Median :19.05 Median :391.44 Median :11.36 Median :21.20
Mean   :18.46 Mean   :356.67 Mean :12.65 Mean :22.53
3rd Qu.:20.20 3rd Qu.:396.23 3rd Qu.:16.95 3rd Qu.:25.00
Max.   :22.00 Max.   :396.90 Max. :37.97 Max. :50.00
```

We see that the crime rate (*crim*) ranges from 0.00632 to 88.97 per capita, the full-value property-tax rate (*tax*) ranges from 187 to 711 per \$10,000, and the pupil-teacher ratio (*ptratio*) ranges from 12.6 to 22. Of these three predictors, the pupil-teacher ratio seems to have a particularly small range, with even the most crowded

classes not exceeding 22 students per teacher on average. This may be due to good state requirements for class sizes. Meanwhile, the property tax rate and especially the crime rate appear to vary drastically. The very large range of the crime rate suggests that some suburbs are clearly more dangerous than others.

```
> which.max(Boston$crim)
[1] 381
> which.max(Boston$tax)
[1] 489
> which.max(Boston$ptratio)
[1] 355
```

We find that suburb #381 has the highest crime rate, suburb #489 has the highest tax rate, and suburb #355 has the highest pupil-teacher ratio.

Part e)

```
> sum(Boston$chas==1)
[1] 35
```

We find that 35 of the suburbs in this data set are bound by the Charles river.

Part f)

```
> summary(Boston$ptratio)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
 12.60   17.40   19.05   18.46   20.20   22.00
```

The median pupil-teacher ratio among the towns in this dataset is 19.05.

Part g)

```
> which.min(Boston$medv)
[1] 399
```

The suburb with the lowest median value of owner-occupied homes is suburb #399.

The values of the other predictors for suburb #399 are shown below:

```
> Boston[399,]
   crim zn indus chas   nox    rm age    dis rad tax ptratio black lstat medv
399 38.3518  0  18.1    0 0.693 5.453 100 1.4896  24  666    20.2 396.9 30.59    5
```

Compared with the overall ranges of the predictors (shown below), suburb #399 appears to have a far above average crime rate, a higher than average proportion of non-retail business acres per town, a greater than average nitrogen oxide concentration, lower than average number of rooms per dwelling, the highest

proportion of owner-occupied units built prior to 1940, the highest proportion of blacks, and a much higher proportion of lower-status individuals. These predictor values make sense since suburb #399 appears to be a much poorer suburb compared to most other Boston suburbs, and thus the predictors general associated with lower socioeconomic situations are respectively higher for this suburb.

```
> summary(Boston)
      crim      zn      indus      chas      nox      rm      age      dis      rad      tax
Min.   :0.00632  Min.   : 0.00  Min.   : 0.46  Min.   :0.00000  Min.   :0.3850  Min.   :3.561  Min.   : 2.90  Min.   : 1.130  Min.   : 1.000  Min.   :187.0
1st Qu.: 0.08294  1st Qu.: 0.00  1st Qu.: 5.19  1st Qu.:0.00000  1st Qu.:0.4490  1st Qu.:5.886  1st Qu.: 45.02  1st Qu.: 2.100  1st Qu.: 4.000  1st Qu.:279.0
Median : 0.25651  Median : 0.00  Median : 9.69  Median :0.00000  Median :0.5380  Median :6.208  Median : 77.50  Median : 3.207  Median : 5.000  Median :330.0
Mean   : 3.61352  Mean   :11.36  Mean   :11.14  Mean   :0.06917  Mean   :0.5547  Mean   :6.285  Mean   : 68.57  Mean   : 3.795  Mean   : 9.549  Mean   :408.2
3rd Qu.: 3.67708  3rd Qu.:12.50  3rd Qu.:18.10  3rd Qu.:0.00000  3rd Qu.:0.6240  3rd Qu.:6.623  3rd Qu.: 94.08  3rd Qu.: 5.188  3rd Qu.:24.000  3rd Qu.:666.0
Max.   :88.97620  Max.   :100.00  Max.   :27.74  Max.   :1.00000  Max.   :0.8710  Max.   :8.780  Max.   :100.00  Max.   :12.127  Max.   :24.000  Max.   :711.0

      ptratio      black      lstat      medv
Min.   :12.60  Min.   : 0.32  Min.   : 1.73  Min.   : 5.00
1st Qu.:17.40  1st Qu.:375.38  1st Qu.: 6.95  1st Qu.:17.02
Median :19.05  Median :391.44  Median :11.36  Median :21.20
Mean   :18.46  Mean   :356.67  Mean   :12.65  Mean   :22.53
3rd Qu.:20.20  3rd Qu.:396.23  3rd Qu.:16.95  3rd Qu.:25.00
Max.   :22.00  Max.   :396.90  Max.   :37.97  Max.   :50.00
```

Part h)


```
> sum(Boston$rm>7)
[1] 64
> sum(Boston$rm>8)
[1] 13
```

We find that 64 suburbs average more than 7 rooms per dwelling, and 13 suburbs average more than 8 rooms per dwelling. Of the suburbs in this latter group (see below), we find that they generally exhibit predictor scores characteristic of more affluent neighborhoods, such as generally having higher than average median values of owner-occupied homes and lower than average proportion of lower status individuals. This suggests that these neighborhoods are on average more affluent and socioeconomically well-off than the average Boston suburb.

```
> which(Boston$rm>8)
[1] 98 164 205 225 226 227 233 234 254 258 263 268 365
> Boston[c(98,164,205,225,226,227,233,234,254,258,263,268,365),]
      crim zn indus chas      nox      rm      age      dis      rad      tax ptratio      black lstat medv
98  0.12083  0  2.89    0 0.4450  8.069  76.0  3.4952    2  276    18.0 396.90  4.21 38.7
164 1.51902  0 19.58    1 0.6050  8.375  93.9  2.1620    5  403    14.7 388.45  3.32 50.0
205 0.02009 95  2.68    0 0.4161  8.034  31.9  5.1180    4  224    14.7 390.55  2.88 50.0
225 0.31533  0  6.20    0 0.5040  8.266  78.3  2.8944    8  307    17.4 385.05  4.14 44.8
226 0.52693  0  6.20    0 0.5040  8.725  83.0  2.8944    8  307    17.4 382.00  4.63 50.0
227 0.38214  0  6.20    0 0.5040  8.040  86.5  3.2157    8  307    17.4 387.38  3.13 37.6
233 0.57529  0  6.20    0 0.5070  8.337  73.3  3.8384    8  307    17.4 385.91  2.47 41.7
234 0.33147  0  6.20    0 0.5070  8.247  70.4  3.6519    8  307    17.4 378.95  3.95 48.3
254 0.36894 22  5.86    0 0.4310  8.259   8.4  8.9067    7  330    19.1 396.90  3.54 42.8
258 0.61154 20  3.97    0 0.6470  8.704  86.9  1.8010    5  264    13.0 389.70  5.12 50.0
263 0.52014 20  3.97    0 0.6470  8.398  91.5  2.2885    5  264    13.0 386.86  5.91 48.8
268 0.57834 20  3.97    0 0.5750  8.297  67.0  2.4216    5  264    13.0 384.54  7.44 50.0
365 3.47428  0 18.10    1 0.7180  8.780  82.9  1.9047   24  666    20.2 354.55  5.29 21.9
```

Problem #4

Part a)




This is the picture of the face for face_01.

```
> class(face_01)
[1] "pixmapGrey"
attr(,"package")
[1] "pixmap"
> face_01@size
[1] 192 168
```

We find that face_01 has the class of pixmapGrey (meaning it is a greyscale image).
We find that the size of face_01 is 192 (height) by 168 (width) pixels.

Part b)



This is the picture of face_01 and face_02 side by side.

```
> range(faces_matrix)
[1] 0.007843137 1.000000000
```

We find that the range of pixel values for these two images ranges from 0.0078 to 1.0. In general, pixels for the class of objects `pixmapGrey` can range from 0 (minimum) to 1 (maximum), where 0 corresponds to the color black and 1 corresponds to the color white.

Part c)

The `dir_list_1` and `dir_list_2` lists capture some of the folder structure of the `CroppedYale` folder. `dir_list_1` captures the first layer of folders under `CroppedYale` (such as `yaleB01`, `yaleB02`, etc.) and these correspond to different people that were photographed. `dir_list_2` captures the files under each subject's folder in the first layer, and these correspond to the different lighting conditions used when photographing each subject.

```
> length(dir_list_1)
[1] 38
> length(dir_list_2)
[1] 2547
```

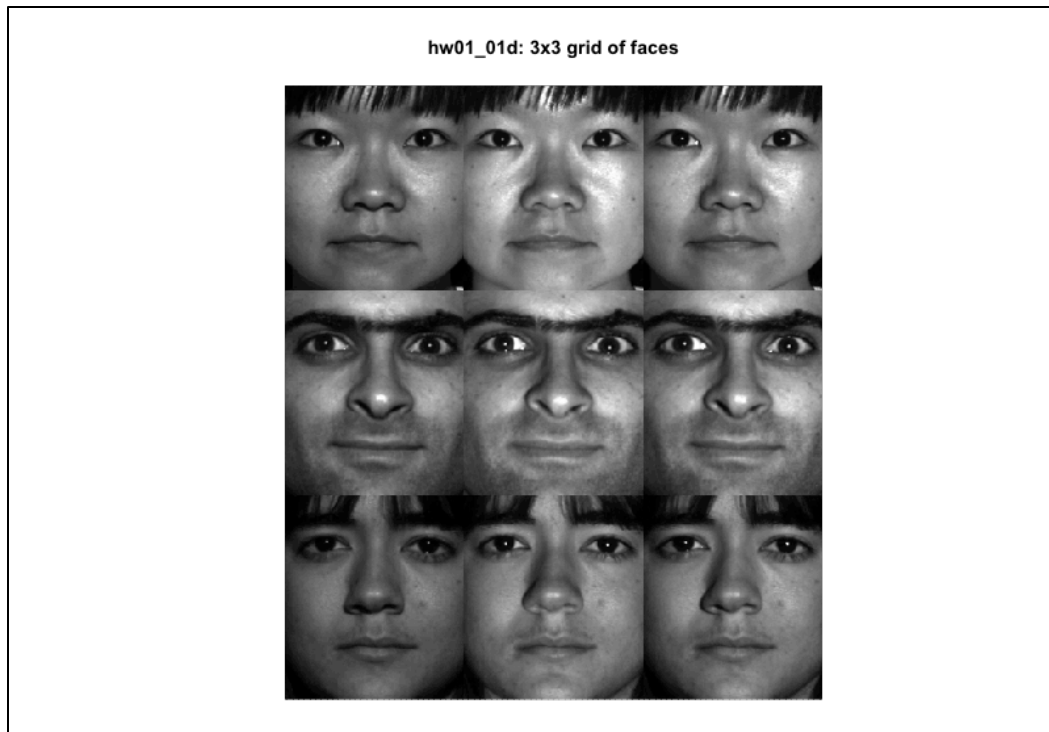
There are 38 elements in `dir_list_1` (meaning 38 subjects photographed) and 2547 elements in `dir_list_2` (meaning 2547 photos taken).

Some sample elements are shown below:

```
> head(dir_list_1)
[1] "yaleB01" "yaleB02" "yaleB03" "yaleB04" "yaleB05" "yaleB06"
> head(dir_list_2)
[1] "yaleB01/DEADJOE" "yaleB01/WS_FTP.LOG" "yaleB01/yaleB01_P00_Ambient.pgm" "yaleB01/yaleB01_P00.info"
[5] "yaleB01/yaleB01_P00A-005E-10.pgm" "yaleB01/yaleB01_P00A-005E+10.pgm"
```

Part d)

Using nested for loops, `rbind` (binding rows), and `cbind` (binding columns), we create a matrix of our images and plot them to generate the following images:



Code

```
#####  
# Kevin Gong  
# STAT W4240  
# Homework 1 , Problem 1  
# 2/5/14  
#  
  
#####  
  
#####  
# Setup  
#####  
  
# make sure R is in the proper working directory  
# note that this will be a different path for every machine  
setwd("~/Dropbox/SIPA/Data Mining")  
  
# first include the relevant libraries  
# note that a loading error might mean that you have to  
# install the package into your R distribution.  
install.packages("ISLR")  
library(ISLR)  
  
#####  
# Problem 1a  
#####  
  
#importing the data  
college <- read.csv('http://www-bcf.usc.edu/~gareth/ISL/College.csv')  
college  
head(college)  
  
#####  
# Problem 1b  
#####  
  
#adding a column with university names  
fix(college)  
rownames(college) <- college[,1]  
fix(college)  
  
#eliminating the original column of university names  
college <- college[,-1]  
fix(college)  
  
#####  
# Problem 1c  
#####  
  
#  
# Part i  
#  
  
#numerical summary of the variables in the data set  
summary(college)
```

```

#
# Part ii
#

#scatterplot matrix of first 10 variables
pairs(college[,1:10])

#
# Part iii
#

#side-by-side boxplots of Private and Outstate
plot(college$Private,college$Outstate)

#
# Part iv
#

#generate new variable Elite based on whether more than 50% of incoming class come from
the top 10% of their high school
Elite=rep("No",nrow(college ))
Elite[college$Top10perc >50]=" Yes"
Elite=as.factor(Elite)
college=data.frame(college ,Elite)

#we see there are 78 Elite universities
summary(college)

#side-by-side boxplots of Private and Outstate
plot(college$Elite,college$Outstate)

#
# Part v
#

par(mfrow=c(2,2))
hist(college$Apps, breaks=5)
hist(college$Apps, breaks=10)
hist(college$Apps, breaks=25)
hist(college$Apps, breaks=75)

par(mfrow=c(2,2))
hist(college$Accept, breaks=5)
hist(college$Accept, breaks=10)
hist(college$Accept, breaks=25)
hist(college$Accept, breaks=75)

par(mfrow=c(2,2))
hist(college$F.Undergrad, breaks=5)
hist(college$F.Undergrad, breaks=10)
hist(college$F.Undergrad, breaks=25)
hist(college$F.Undergrad, breaks=75)

par(mfrow=c(2,2))

```

```

hist(college$Expend, breaks=5)
hist(college$Expend, breaks=10)
hist(college$Expend, breaks=25)
hist(college$Expend, breaks=75)

par(mfrow=c(2,2))
hist(college$Grad.Rate, breaks=5)
hist(college$Grad.Rate, breaks=10)
hist(college$Grad.Rate, breaks=25)
hist(college$Grad.Rate, breaks=75)

#
# Part vi
#

mean(college$Expend[college$Elite=="Yes"])
summary(college$Expend,college$Elite=="No")

```

```

#####
# Kevin Gong
# STAT W4240
# Homework 1 , Problem 2
# 2/5/14
#
#####

```

```

#####
# Setup
#####

```

```

# make sure R is in the proper working directory
# note that this will be a different path for every machine
setwd("~/Dropbox/SIPA/Data Mining")

# first include the relevant libraries
# note that a loading error might mean that you have to
# install the package into your R distribution.
install.packages("ISLR")
library(ISLR)

```

```

#####
# Problem 1a
#####

```

```

#importing the data
auto <- read.csv("~/Dropbox/SIPA/Data Mining/Autodata3.csv")

```

```

#examine which ones are quantitative/qualitative
summary(auto)
sapply(auto,class)

```

```

#quantitative: mpg, cylinders, displacement, horsepower, weight, acceleration, year,
  origin
#qualitative: name

```

```
#####  
# Problem 1b  
#####
```

```
#range of each quantitative predictor  
sapply(auto[,1:8],range,na.rm=TRUE)
```

```
#####  
# Problem 1c  
#####
```

```
#mean and standard deviation of each quantitative predictor  
sapply(auto[,1:8],mean,na.rm=TRUE)  
sapply(auto[,1:8],sd,na.rm=TRUE)
```

```
#####  
# Problem 1d  
#####
```

```
#removing rows 10 through 85  
auto = auto[-c(10:85),]
```

```
#range, mean, and standard deviation of remaining observations of each quantitative  
predictor  
sapply(auto[,1:8],range,na.rm=TRUE)  
sapply(auto[,1:8],mean,na.rm=TRUE)  
sapply(auto[,1:8],sd,na.rm=TRUE)
```

```
#####  
# Problem 1e  
#####
```

```
#restore removed rows  
auto <- read.csv('~/.Dropbox/SIPA/Data Mining/Autodata3.csv')
```

```
#comparing relationships between predictors via scatterplots  
pairs(auto)
```

```
#####  
# Kevin Gong  
# STAT W4240  
# Homework 1 , Problem 3  
# 2/5/14  
#  
# The following code loads the eigenfaces data and  
# performs a set of simple loading and plotting functions  
#####
```

```
#####
```



```

# Setup
#####

# make sure R is in the proper working directory
# note that this will be a different path for every machine
setwd("~/Dropbox/SIPA/Data Mining")

# first include the relevant libraries
# note that a loading error might mean that you have to
# install the package into your R distribution.
install.packages("MASS")
library(MASS)

#####
# Problem 3a
#####

#count the number of rows in Boston
?Boston

#####
# Problem 3b
#####

#creating pairwise scatterplots
pairs(Boston)

#####
# Problem 3d
#####

#range, mean, and standard deviation of Boston predictors
sapply(Boston,range,na.rm=TRUE)
sapply(Boston,mean,na.rm=TRUE)
sapply(Boston,sd,na.rm=TRUE)
summary(Boston)

#find which have highest crime rates, tax rates, and pupil-teacher ratios
which.max(Boston$crim)
which.max(Boston$tax)
which.max(Boston$ptratio)

#####
# Problem 3e
#####

#counting the number of suburbs bordering the Charles River
sum(Boston$chas==1)
sum(Boston$chas==0)
sum(is.na(Boston$chas)) #check for missing data

#####
# Problem 3f
#####

```

```

#median pupil-teacher teacher among the towns
summary(Boston$ptratio)

#####
# Problem 3g
#####

#finding the lowest median value of owner-occupied homes
which.min(Boston$medv)

#values of the other predictors for suburb 399
Boston[399,]

#compare the predictor values of suburb 399 with overall predictor ranges
summary(Boston)

#####
# Problem 3h
#####

#finding the number of suburbs averaging more than 7 rooms per dwelling
sum(Boston$rm>7)

#finding the number of suburbs averaging more than 8 rooms per dwelling
sum(Boston$rm>8)

#compare the predictor values of the suburb averaging more than 8 rooms per dwelling
which(Boston$rm>8)
Boston[c(98,164,205,225,226,227,233,234,254,258,263,268,365),]
summary(Boston)


#####
# Kevin Gong
# STAT W4240
# Homework 1 , Problem 4
# 2/5/14
#
# The following code loads the eigenfaces data and
# performs a set of simple loading and plotting functions
#####

#####
# Setup
#####

# make sure R is in the proper working directory
# note that this will be a different path for every machine
setwd("~/Dropbox/SIPA/Data Mining")

# first include the relevant libraries

```

```

# note that a loading error might mean that you have to
# install the package into your R distribution. From the
# command line, type install.packages("pixmap")
library(pixmap)

#####
# Problem 1a
#####

# paste or type in the given code here
face_01 = read.pnm(file = "CroppedYale/yaleB01/yaleB01_P00A-005E+10.pgm")

# now plot the data
plot(face_01)
# give it a nice title
title('hw01_01a: the first face')
# save the result
filename = 'hw01_01a.png'
dev.copy(device=png, file=filename, height=600, width=800)
dev.off()

# extract the class and size

#----- START YOUR CODE BLOCK HERE -----#
class(face_01)
face_01@size
#alternative method to get size
nrow(getChannels(face_01))
ncol(getChannels(face_01))
#----- END YOUR CODE BLOCK HERE -----#

#####
# Problem 1b
#####

# make face_01 into a matrix with the given command
face_01_matrix = getChannels(face_01)

# load a second face
face_02 = read.pnm(file = "CroppedYale/yaleB02/yaleB02_P00A-005E+10.pgm")
face_02_matrix = getChannels(face_02)

# combine two faces into a single data matrix and make that a pixmap
faces_matrix = cbind( face_01_matrix , face_02_matrix )
faces = pixmapGrey( faces_matrix )

# plot to verify
plot(faces)

# find min and max values

#----- START YOUR CODE BLOCK HERE -----#
range(faces_matrix)
#alternative method to get min and max values
min(faces_matrix)
max(faces_matrix)
#----- END YOUR CODE BLOCK HERE -----#

#####
# Problem 1c
#####

```

```

# get directory structure
dir_list_1 = dir(path="CroppedYale/",all.files=FALSE)
dir_list_2 = dir(path="CroppedYale/",all.files=FALSE,recursive=TRUE)

# find lengths

#----- START YOUR CODE BLOCK HERE -----#
length(dir_list_1)
length(dir_list_2)
#example elements
head(dir_list_1)
head(dir_list_2)
#----- END YOUR CODE BLOCK HERE -----#

#####
# Problem 1d
#####

# the list of pictures (note the absence of 14 means that 31 corresponds to yaleB32)
pic_list = c( 05 , 11 , 31 )
view_list = c( 'P00A-005E+10' , 'P00A-005E-10' , 'P00A-010E+00' )

# preallocate an empty list
pic_data = vector("list",length(pic_list)*length(view_list))
# initialize an empty matrix of faces data
faces_matrix = vector()

#----- START YOUR CODE BLOCK HERE -----#

pic_list = c( 05 , 11 , 31 )
view_list = c( 'P00A-005E+10' , 'P00A-005E-10' , 'P00A-010E+00' )

faces_matrix = vector()
placeholder_matrix= vector()

for (i in 1:3) {

placeholder_matrix <- NULL

for (j in 1:3) {
  filename=
  sprintf("CroppedYale/%s/%s_%s.pgm",dir_list_1[pic_list[i]],dir_list_1[pic_list[i]],view_l
ist[j])
  print(filename)
  face=read.pnm(file=filename)
  pic_data <-getChannels(face)
  placeholder_matrix=cbind(placeholder_matrix,pic_data)
}
faces_matrix=rbind(faces_matrix,placeholder_matrix)
}

#----- END YOUR CODE BLOCK HERE -----#

# now faces_matrix has been built properly. plot and save it.
faces = pixmapGrey(faces_matrix)
plot(faces)
# give it a nice title

```

```
title('hw01_01d: 3x3 grid of faces')
# save the result
filename = 'hw01_01d.png'
dev.copy(device=png, file=filename, height=600, width=800)
dev.off()
```

```
#####
# End of Script
#####
```